

ExampleManipulatingDataFrames

May 12, 2022

In this lecture I'm going to walk through a basic data cleaning process with you and introduce you to a few more pandas API functions.

```
[12]: # Let's start by bringing in pandas
import pandas as pd
# And load our dataset. We're going to be cleaning the list of presidents in
→ the US from wikipedia
df=pd.read_csv("datasets/presidents.csv")
# And lets just take a look at some of the data
df.head()
```

```
[12]: #      #      President      Born      Age atstart of presidency \
0  1  George Washington  Feb 22, 1732[a]  57ăyears, 67ădaysApr 30, 1789
1  2      John Adams    Oct 30, 1735[a]  61ăyears, 125ădaysMar 4, 1797
2  3  Thomas Jefferson  Apr 13, 1743[a]  57ăyears, 325ădaysMar 4, 1801
3  4      James Madison  Mar 16, 1751[a]  57ăyears, 353ădaysMar 4, 1809
4  5      James Monroe   Apr 28, 1758   58ăyears, 310ădaysMar 4, 1817
```

```
      Age atend of presidency Post-presidencytimespan      Died \
0  65ăyears, 10ădaysMar 4, 1797      2ăyears, 285ădays  Dec 14, 1799
1  65ăyears, 125ădaysMar 4, 1801      25ăyears, 122ădays   Jul 4, 1826
2  65ăyears, 325ădaysMar 4, 1809      17ăyears, 122ădays   Jul 4, 1826
3  65ăyears, 353ădaysMar 4, 1817      19ăyears, 116ădays  Jun 28, 1836
4  66ăyears, 310ădaysMar 4, 1825       6ăyears, 122ădays   Jul 4, 1831
```

```
      Age
0  67ăyears, 295ădays
1  90ăyears, 247ădays
2  83ăyears, 82ădays
3  85ăyears, 104ădays
4  73ăyears, 67ădays
```

```
[19]: # Ok, we have some presidents, some dates, I see a bunch of footnotes in the
→ "Born" column which might cause
# issues. Let's start with cleaning up that name into firstname and lastname.
→ I'm going to tackle this with
# a regex. So I want to create two new columns and apply a regex to the
→ projection of the "President" column.
```

```
# Here's one solution, we could make a copy of the President column
df["First"]=df['President']

# Then we can call replace() and just have a pattern that matches the last name
→and set it to an empty string

df["First"]=df["First"].replace("[ ].*", "", regex=True)
# buscamos en la columna "First" un patrón que contenga un espacio([ ]) seguido
→de numerosos caracteres(.*?) y que lo reemplace por un string
# vacío ("") en la misma columna donde lo buscamos

# Now let's take a look
df.head()
```

```
[19]: #      President      Born      Age atstart of presidency \
0 1 George Washington Feb 22, 1732[a] 57ăyears, 67ădaysApr 30, 1789
1 2      John Adams Oct 30, 1735[a] 61ăyears, 125ădaysMar 4, 1797
2 3 Thomas Jefferson Apr 13, 1743[a] 57ăyears, 325ădaysMar 4, 1801
3 4      James Madison Mar 16, 1751[a] 57ăyears, 353ădaysMar 4, 1809
4 5      James Monroe Apr 28, 1758 58ăyears, 310ădaysMar 4, 1817
```

```
      Age atend of presidency Post-presidencytimespan      Died \
0 65ăyears, 10ădaysMar 4, 1797      2ăyears, 285ădays Dec 14, 1799
1 65ăyears, 125ădaysMar 4, 1801      25ăyears, 122ădays Jul 4, 1826
2 65ăyears, 325ădaysMar 4, 1809      17ăyears, 122ădays Jul 4, 1826
3 65ăyears, 353ădaysMar 4, 1817      19ăyears, 116ădays Jun 28, 1836
4 66ăyears, 310ădaysMar 4, 1825      6ăyears, 122ădays Jul 4, 1831
```

```
      Age First
0 67ăyears, 295ădays George
1 90ăyears, 247ădays John
2 83ăyears, 82ădays Thomas
3 85ăyears, 104ădays James
4 73ăyears, 67ădays James
```

```
[20]: # That works, but it's kind of gross. And it's slow, since we had to make a
→full copy of a column then go
# through and update strings. There are a few other ways we can deal with this.
→Let me show you the most
# general one first, and that's called the apply() function. Let's drop the
→column we made first
del(df["First"])

# The apply() function on a dataframe will take some arbitrary function you
→have written and apply it to
```

```

# either a Series (a single column) or DataFrame across all rows or columns.
→ Lets write a function which
# just splits a string into two pieces using a single row of data
def splitname(row):
    # The row is a single Series object which is a single row indexed by column
    → values
    # Let's extract the firstname and create a new entry in the series
    row['First']=row['President'].split(" ")[0]
    # Let's do the same with the last word in the string
    row['Last']=row['President'].split(" ")[-1]
    # Now we just return the row and the pandas .apply() will take of merging
    → them back into a DataFrame
    return row

# Now if we apply this to the dataframe indicating we want to apply it across
→ columns
df=df.apply(splitname, axis='columns')
df.head()

```

[20]:

#	President	Born	Age atstart of presidency \
0	1 George Washington	Feb 22, 1732[a]	57ăyears, 67ădaysApr 30, 1789
1	2 John Adams	Oct 30, 1735[a]	61ăyears, 125ădaysMar 4, 1797
2	3 Thomas Jefferson	Apr 13, 1743[a]	57ăyears, 325ădaysMar 4, 1801
3	4 James Madison	Mar 16, 1751[a]	57ăyears, 353ădaysMar 4, 1809
4	5 James Monroe	Apr 28, 1758	58ăyears, 310ădaysMar 4, 1817

	Age atend of presidency	Post-presidencytimespan	Died \
0	65ăyears, 10ădaysMar 4, 1797	2ăyears, 285ădays	Dec 14, 1799
1	65ăyears, 125ădaysMar 4, 1801	25ăyears, 122ădays	Jul 4, 1826
2	65ăyears, 325ădaysMar 4, 1809	17ăyears, 122ădays	Jul 4, 1826
3	65ăyears, 353ădaysMar 4, 1817	19ăyears, 116ădays	Jun 28, 1836
4	66ăyears, 310ădaysMar 4, 1825	6ăyears, 122ădays	Jul 4, 1831

	Age	First	Last
0	67ăyears, 295ădays	George	Washington
1	90ăyears, 247ădays	John	Adams
2	83ăyears, 82ădays	Thomas	Jefferson
3	85ăyears, 104ădays	James	Madison
4	73ăyears, 67ădays	James	Monroe

[21]:

```

# Pretty questionable as to whether that is less gross, but it achieves the
→ result and I find that I use the
# apply() function regularly in my work. The pandas series has a couple of
→ other nice convenience functions
# though, and the next I would like to touch on is called .extract(). Lets drop
→ our firstname and lastname.
del(df['First'])

```

```
del(df['Last'])

# Extract takes a regular expression as input and specifically requires you to
→ set capture groups that
# correspond to the output columns you are interested in. And, this is a great
→ place for you to pause the
# video and reflect - if you were going to write a regular expression that
→ returned groups and just had the
# firstname and lastname in it, what would that look like?

# Here's my solution, where we match three groups but only return two, the
→ first and the last name
pattern="(^[\w]*)(?:.* )([\w]*$)"
# Acá el segundo grupo, al poner ?: se indica que ese grupo NO debe agregarse
→ al resultado devuelto por REGEX
# Por eso el resultado es una matriz con 2 columnas, Nombre (grupo1) y Apellido
→ (grupo2)

# Now the extract function is built into the str attribute of the Series
→ object, so we can call it
# using Series.str.extract(pattern)
df["President"].str.extract(pattern).head()
```

```
[21]:      0      1
0  George Washington
1    John      Adams
2  Thomas  Jefferson
3   James    Madison
4   James    Monroe
```

```
[5]: # So that looks pretty nice, other than the column names. But if we name the
→ groups we get named columns out

# acá es lo mismo que arriba, mismo patrón pero a los grupos 1 y 3 se le
→ asignan nombres para que aparezcan en las columnas
pattern="(P<First>^[\\w]*)(?:.* )(P<Last>[\\w]*$)"

# Now call extract

# Se crea una nueva matriz a partir de la columna de la matriz original df
names=df["President"].str.extract(pattern).head()
names
```

```
[5]:      First      Last
0  George Washington
1    John      Adams
```

```

2 Thomas Jefferson
3 James Madison
4 James Monroe

```

[6]: *# And we can just copy these into our main dataframe if we want to*

```

# Copiamos las nuevas columnas dentro de la matriz original
df["First"]=names["First"]
df["Last"]=names["Last"]
df.head()

```

[6]:

	#	President	Born	Age atstart of presidency \
0	1	George Washington	Feb 22, 1732[a]	57ăyears, 67ădaysApr 30, 1789
1	2	John Adams	Oct 30, 1735[a]	61ăyears, 125ădaysMar 4, 1797
2	3	Thomas Jefferson	Apr 13, 1743[a]	57ăyears, 325ădaysMar 4, 1801
3	4	James Madison	Mar 16, 1751[a]	57ăyears, 353ădaysMar 4, 1809
4	5	James Monroe	Apr 28, 1758	58ăyears, 310ădaysMar 4, 1817

	Age atend of presidency	Post-presidencytimespan	Died \
0	65ăyears, 10ădaysMar 4, 1797	2ăyears, 285ădays	Dec 14, 1799
1	65ăyears, 125ădaysMar 4, 1801	25ăyears, 122ădays	Jul 4, 1826
2	65ăyears, 325ădaysMar 4, 1809	17ăyears, 122ădays	Jul 4, 1826
3	65ăyears, 353ădaysMar 4, 1817	19ăyears, 116ădays	Jun 28, 1836
4	66ăyears, 310ădaysMar 4, 1825	6ăyears, 122ădays	Jul 4, 1831

	Age	First	Last
0	67ăyears, 295ădays	George	Washington
1	90ăyears, 247ădays	John	Adams
2	83ăyears, 82ădays	Thomas	Jefferson
3	85ăyears, 104ădays	James	Madison
4	73ăyears, 67ădays	James	Monroe

[7]: *# It's worth looking at the pandas str module for other functions which have*
→been written specifically

to clean up strings in DataFrames, and you can find that in the docs in the
→Working with Text

section: https://pandas.pydata.org/pandas-docs/stable/user_guide/text.html

[8]: *# Now lets move on to clean up that Born column. First, let's get rid of*
→anything that isn't in the

pattern of Month Day and Year

usamos regex para extraer únicamente las fechas compuestas por MMM DD, AAAA
se reemplaza la columna Born por los nuevos datos

```

df["Born"]=df["Born"].str.extract("([\w]{3} [\w]{1,2}, [\w]{4})")
df["Born"].head()

```

```
[8]: 0    Feb 22, 1732
      1    Oct 30, 1735
      2    Apr 13, 1743
      3    Mar 16, 1751
      4    Apr 28, 1758
      Name: Born, dtype: object
```

```
[9]: # So, that cleans up the date format. But I'm going to foreshadow something
      →else here - the type of this
      # column is object, and we know that's what pandas uses when it is dealing with
      →string. But pandas actually
      # has really interesting date/time features - in fact, that's one of the
      →reasons Wes McKinney put his efforts
      # into the library, to deal with financial transactions. So if I were building
      →this out, I would actually
      # update this column to the write data type as well
      df["Born"]=pd.to_datetime(df["Born"])
      df["Born"].head()
```

```
[9]: 0    1732-02-22
      1    1735-10-30
      2    1743-04-13
      3    1751-03-16
      4    1758-04-28
      Name: Born, dtype: datetime64[ns]
```

```
[10]: # This would make subsequent processing on the dataframe around dates, such as
      →getting every President who
      # was born in a given time span, much easier.
```

Now, most of the other columns in this dataset I would clean in a similar fashion. And this would be a good practice activity for you, so I would recommend that you pause the video, open up the notebook for the lecture if you don't already have it opened, and then finish cleaning up this dataframe. In this lecture I introduced you to the `str` module which has a number of important functions for cleaning pandas dataframes. You don't have to use these - I actually use `apply()` quite a bit myself, especially if I don't need high performance data cleaning because my dataset is small. But the `str` functions are incredibly useful and build on your existing knowledge of regular expressions, and because they are vectorized they are efficient to use as well.